

Search...

[UNIX Tutorials for Beginners](#)

[Introduction to Unix](#)

[Tutorial 1: Files in Unix](#)

[Tutorial 2: Working with Files](#)

[Tutorial 3: More About Files](#)

[Tutorial 4: Wildcards and File Access](#)

[Tutorial 5: Processes and Jobs](#)

[Tutorial 6: Other Useful Commands](#)

[Tutorial 7: Environment Variables](#)

# UNIX Tutorials for Beginners

These tutorials are derived from the excellent tutorials from the University of Surrey, UK, with some minor modifications for our site. The originals can be found [here](#).

Last updated on Apr 29, 2019



© 2023 The Rector and Visitors of the University of Virginia





Search...

UNIX Tutorials for Beginners

Introduction to Unix

Tutorial 1: Files in Unix

Tutorial 2: Working with Files

Tutorial 3: More About Files

Tutorial 4: Wildcards and File Access

Tutorial 5: Processes and Jobs

Tutorial 6: Other Useful Commands

Tutorial 7: Environment Variables

# Introduction to Unix

## Typographical Conventions

In what follows, we shall use the following typographical conventions:

- Characters written in color-coded typewriter font are commands to be typed into the computer as they stand.
- Characters written in standard typewriter font indicate non-specific file or directory names.
- Words inserted within square brackets, e.g. [Ctrl], indicate keys to be pressed.

So, for example:

```
% ls anydirectory [Enter]
```

means “at the Unix prompt %, type `ls` followed by the name of some directory, then press the key marked [Enter].” Don’t forget to press the [Enter] key: commands are not sent to the computer until this is done.

Note: Unix is case-sensitive, so “LS” is not the same as `ls`. The same applies to filenames, so `myfile.txt`, `MyFile.txt` and `MYFILE.TXT` are three separate files. Beware if copying files to a PC, since DOS and Windows do not make this distinction.

## Introduction to the Unix Operating System

An *operating system* is the suite of programs that make the computer work. Historically, Unix had two “flavors,” System V from AT&T’s Bell Labs, and BSD (for Berkeley Software Distribution). Today that distinction is less relevant but the dominant versions of Unix still have different roots. *Linux*, used for some workstations and many high-performance computing clusters, arose from System V. Mac OS was derived from a version of BSD. Android is a heavily-modified Linux.

The “windowing” (graphical) system on Linux is called X11. Modern Linux systems all provide at least one “desktop” environment, similar to Windows or the Mac OS desktop. However, working at the command line is still often the most efficient use of teh system, so users can benefit by learning their way around a terminal.

A Unix operating system is made up of three parts: the kernel, the shell and the programs (applications or “apps”).

### The Kernel

The kernel of Unix is the hub of the operating system: it allocates time and memory to programs and handles the filesystem and communications in response to system calls.

### The Shell

The shell acts as an interface between the user and the kernel. When a user logs in, the login program checks the username and password, and then starts another program called the shell. The shell is a command line interpreter (CLI). It interprets the commands the user types in and arranges for them to be carried out. The commands are themselves programs: when they terminate, the shell gives the user another prompt. The adept user can customise his/her own shell, and users can use different shells on the same machine. The `bash` shell is the default on Linux. Mac OS formerly used bash, but on newer releases the `zsh` shell is the default. Most of the basic commands are the same.

As an illustration of the way that the shell and the kernel work together, suppose a user types `rm myfile` to delete the file `myfile`. The shell searches the filesystem, first for the file containing the program `rm`, and directs the kernel, through system calls, to *execute* the program `rm` on `myfile`. When the process `rm myfile` has finished, the shell then returns the prompt to the user, indicating that it is ready for further commands.

### Some Features

- Filename Tab Completion: By typing part of the name of a command, filename or directory and pressing the [Tab] key, the shell will complete the rest of the name automatically. If the shell finds more than one name beginning with those letters you have typed, it will halt, prompting you to type a few more letters before pressing the [Tab] key again.
  - Both bash and zsh support tab completion, but zsh has a somewhat more sophisticated set of features.
- History: The shell keeps a list of the commands you have typed in. If you need to repeat a command, use the cursor keys to scroll up and down the list or type `history` for a list of previous commands.
  - Several discussions of the differences between bash and zsh are available online, such as [here](#). The rest of these tutorials are targeted to bash.

### Prompt

The prompt is a character or string of characters that indicates the shell is ready for a command. It varies by shell and system, and can be customized by more advanced users. Throughout these tutorials we will use the percent sign `%` as the prompt.

### NEXT

[Files in Unix](#)

---

Last updated on Jan 1, 0001



Research Computing

© 2023 The Rector and Visitors of the University of Virginia



## UNIX Tutorials for Beginners

### Introduction to Unix

#### [Tutorial 1: Files in Unix](#)

#### [Tutorial 2: Working with Files](#)

#### [Tutorial 3: More About Files](#)

#### [Tutorial 4: Wildcards and File Access](#)

#### [Tutorial 5: Processes and Jobs](#)

#### [Tutorial 6: Other Useful Commands](#)

#### [Tutorial 7: Environment Variables](#)

# Files in Unix

## Files and Processes

Everything in Unix is either a file or a process. A process is an executing program identified by a unique PID (process identifier). A file is a collection of data. They are created by users using text editors, running compilers etc. Examples of files:

- a document (report, essay etc.)
- the text of a program written in some high-level programming language
- instructions comprehensible directly to the machine and incomprehensible to a casual user, for example, a collection of binary digits (an executable or binary file)
- a directory, containing information about its contents, which may be a mixture of other directories (subdirectories) and ordinary files

## The Directory Structure

In Unix, folders are generally called *directories*. Directories are arranged in a hierarchical structure, like an inverted tree. The top of the hierarchy is traditionally called **root**.

## Listing Files and Directories

### **ls** (list)

When you first log in, your current working directory is your *home directory*. Your home directory has the same name as your username, for example, **mst3k**, and it is where your personal files and subdirectories are saved. To find out what is in your home directory, type

```
% ls
```

This is short for “list.” Most Unix commands are two to four letters, sometimes with nonintuitive meanings.

The **ls** command lists the contents of your current working directory. There may be no files visible in your home directory, in which case the prompt will be returned. Alternatively, there may already be some files or folders created when your account was set up.

In most Unix systems, **ls** does not report *hidden files* by default. Files or directories with names beginning with a dot (.) are hidden and usually contain important program configuration information. They are hidden because you should not change them unless you understand what you are doing. To list all files in your home directory including those whose names begin with a dot, type

```
% ls -a
```

The `-a` is an example of a *command-line option*. The options change the behavior of the command. There are online manual pages that tell you which options a particular command can take, and how each option modifies the behavior of the command.

## Making Directories

### **mkdir** (make directory)

We will now make a subdirectory in your home directory to hold the files you will be creating and using in the course of this tutorial. To make a subdirectory called `unixstuff` in your current working directory type

```
% mkdir unixstuff
```

To see the directory you have just created, type

```
% ls
```

## Changing to a Different Directory

### cd (change directory)

The command `cd _directory_` means change the current working directory to “directory”. The current working directory may be thought of as the directory you are in, i.e. your current position in the directory tree. To change to the directory you have just made, type

```
% cd unixstuff
```

Type `ls` to see the directory’s contents (it should be empty).

### Exercise 1A

Make another directory inside the unixstuff directory called backups.

## The Directories `.` and `..`

While still in the unixstuff directory, type

```
% ls -a
```

The unixstuff directory (and in all other directories) contains two special directories called `(.)` and `(..)`. In Unix, `(.)` means the current directory, so typing

```
% cd .
```

means stay where you are (the unixstuff directory). This may not seem very useful at first, but using `(.)` as the name of the current directory will save a lot of typing, as we shall see later in the tutorial. `(..)` means the parent of the current directory, so typing

```
% cd ..
```

will take you one directory up the hierarchy. Try it now.

Note: there is a space between `cd` and the dot or double dot. Also note: typing `cd` with no argument always returns you to your home directory. This is very useful if you are lost in the file system.

## Pathnames

### pwd (print working directory)

Pathnames enable you to work out where you are in relation to the whole filesystem. For example, to find out the absolute pathname of your home directory, type `cd` to get back to your home directory and then type

```
% pwd
```

Most of the time you should see `/home/mst3k`. On a multiuser central system like Rivanna, `/home` may be a *symbolic link*, i.e. an alias, for something else. To see it, type

```
% pwd -P
```

The full pathname may look something like this: `/sfs/qumulo/qhome/mst3k`.

### Exercise 1B

Use the commands `ls`, `pwd` and `cd` to explore the file system. (Remember, if you get lost, type `cd` with no argument to return to your home directory.)

# More About Home Directories and Pathnames

## Understanding Paths

First type `cd` to get back to your home directory, then type

```
% ls unixstuff
```

to list the contents of your `unixstuff` directory. Now type

```
% ls backups
```

You will get a message like this:

```
backups: No such file or directory
```

The reason is that “backups” is not in your current working directory. To use a command on a file (or directory) not in the current working directory (the directory you are currently in), you must either `cd` to the correct directory, or specify its full *pathname*. To list the contents of your `backups` directory, you must type

```
% ls unixstuff/backups
```

This path starts from the current location. A path may be *absolute* or *relative*. An absolute path starts from the root location. The absolute path is

```
ls /home/mst3k/unixstuff/backups
```

A relative path starts from the current working directory. The special symbols `.` and `..` are often used for relative paths.

## ~ (your home directory)

Home directories can also be referenced by the tilde `~` character. It can be used to specify paths starting at your home directory. So typing

```
% ls ~/unixstuff
```

will list the contents of your `unixstuff` directory, no matter where you currently are in the file system. What do you think

```
% ls ~
```

would list? What do you think

```
% ls ~/..
```

would list?

## Getting Help

### Online Manuals

There are online manuals which give information about most commands. The manual pages tell you which options a particular command can take, and how each option modifies the behaviour of the command. Type `man` command to read the manual page for a particular command. For example, to find out more about the `wc` (word count) command, type

```
% man wc
```

In newer Linux systems, most commands have a `--help` option

```
% wc --help
```

However, `man` sends output through a [pager](#), whereas `--help` prints directly to the console.

Another useful command,

```
% whatis wc
```

gives a one-line description of the command, but omits any information about options, etc.

## Summary

Command	Operation
ls	list files and directories
ls -a	list all files and directories
mkdir	make a directory
cd directory	change to directory
cd	change to home directory
cd ~	change to home directory
cd ..	change to parent directory
pwd	display the path of the current directory
man	display the manual pages for the specified command
whatis	display a description of the specified command

PREVIOUS

[Introduction to Unix](#)

NEXT

[Working with Files](#)

Last updated on Jan 1, 0001



Search...

[UNIX Tutorials for Beginners](#)

[Introduction to Unix](#)

[Tutorial 1: Files in Unix](#)

[Tutorial 2: Working with Files](#)

[Tutorial 3: More About Files](#)

[Tutorial 4: Wildcards and File Access](#)

[Tutorial 5: Processes and Jobs](#)

[Tutorial 6: Other Useful Commands](#)

[Tutorial 7: Environment Variables](#)

# Working with Files

## Copying Files

### cp (copy)

`cp file1 file2` is the command which makes a copy of `file1` in the current working directory and calls it `file2`.

For our example we will create a file *science.txt*. Click the down-arrow icon to download the file. Use whatever method you know to place this file into your home directory.

```
In the space of one hundred and seventy-six years the Lower Mississippi has shortened itself  
  
- Mark Twain, Life on the Mississippi
```



```
% cd ~/unixstuff
```

Then at the Unix prompt, type,

```
% cp ../science.txt .
```

Note: Don't forget the dot (.) at the end. Remember, in UNIX, the dot means the current directory.

The above command means copy the file `science.txt` from the parent directory to the current directory, keeping the name the same. To change the name, use

```
% cp ../science.txt ./newname
```

### Exercise 2A

Create a backup of your `science.txt` file by copying it to a file called `science.bak`.

## Moving Files

### mv (move)

`mv file1 file2` moves (or renames) `file1` to `file2`. To move a file from one place to another, use the `mv` command. This has the effect of moving rather than copying the file, so you end up with only one file rather than two. It can also be used to rename a file, by moving the file to the same directory, but giving it a different name. We are now going to move the file `science.bak` to your backup directory. First, change directories to your `unixstuff` directory (can you remember how?). Then, inside the `unixstuff` directory, type

```
% mv science.bak backups/.
```

Type `ls` and `ls backups` to see if it has worked.

## Removing Files and Directories

### rm (remove), rmdir (remove directory)

To delete (remove) a file, use the `rm` command. As an example, we are going to create a copy of the `science.txt` file then delete it. Inside your `unixstuff` directory, type

```
% cp science.txt tempfile.txt
```

Confirm the file was created:

```
% ls
```

Now delete it:

```
% rm tempfile.txt
% ls
```

You can use the `rmdir` command to remove a directory, but only if it is empty. Try to remove the backups directory. You will not be able to since Unix will not let you remove a non-empty directory.

To remove a non-empty directory use

```
rm -rf directory
```



The above command will remove the directory without confirming anything! Be extremely careful with it!

You can request confirmation with

```
% rm -if directory
```

though this may be tedious. The `-i` option (inquire) also works for `rm`

```
% rm -i myfile
```

## Exercise 2B

Create a directory called tempstuff using `mkdir`, then remove it using the `rmdir` command.

## Displaying the Contents of a File on the Screen

### cat (concatenate)

The `cat` command can show a text file's contents

```
% cat science.txt
```

Be sure to use the correct path to the file. Cat can also join two text files, hence its name.

```
% cat file1 file2 > file3
```

The `>` sign is a *redirection*, which we will discuss later.

### clear (clear screen)

Clear the screen.

### more

The command `more` prints the contents of a file onto the screen a page at a time. Type

```
% more science.txt
```

Press the [spacebar] if you want to see another page. Type [q] if you want to quit reading.

The `more` command is an example of a **pager**, a program that “pages” through a text file.

### head

The head command writes the first ten lines of a file to the screen.

```
% head science.txt
```

Now type

```
% head -5 science.txt
```

What difference did the `-5` make to the `head` command?

## tail

The `tail` command writes the last ten lines of a file to the screen. Clear the screen and type

```
% tail science.txt
```

How can you view the last 15 lines of the file?

# Searching the Contents of a File

## Simple Searching Using more

Using `more`, you can search through a text file for a keyword (pattern). For example, to search through `science.txt` for the word 'science', type

```
% more science.txt
```

then, still in `more` (i.e. don't press [q] to quit), type a forward slash [/] followed by the word to search

```
/science
```

The `more` command finds and highlights the keyword. Type [n] to search for the next occurrence of the word.

## grep (don't ask why it is called grep)

`grep` is one of many standard Unix utilities. It searches files for specified words or patterns. First clear the screen, then type

```
% grep science science.txt
```

As you can see, `grep` has printed out each line containing the word science... or has it? Try typing

```
% grep Science science.txt
```

The `grep` command is case sensitive; it distinguishes between Science and science. To ignore upper/lower case distinctions, use the `-i` option, i.e. type

```
% grep -i science science.txt
```

To search for a phrase or pattern, you must enclose it in single quotes (the apostrophe symbol). For example to search for spinning top, type

```
% grep -i 'spinning top' science.txt
```

Some of the other options of `grep` are: `-v` (display those lines that do NOT match); `-n` (precede each matching line with the line number); and `-c` (print only the total count of matched lines). Try some of them and see the different results. Don't forget, you can use more than one option at a time, for example, the number of lines without the words `science` or `Science` is

```
% grep -ivc science science.txt
```

## wc (word count)

A handy little utility is the `wc` command, short for word count. To do a word count on `science.txt`, type

```
% wc -w science.txt
```

To find out how many lines the file has, type

```
% wc -l science.txt
```

## Summary

Command	Operation
cp file1 file2	copy file1 and call it file2
mv file1 file2	move or rename file1 to file2
rm file	remove file
rmdir directory	remove directory
cat file	display file
more file	display file a page at a time
head file	display the first few lines of a file
tail file	display the last few lines of file
grep 'keyword' file	search file for keywords
wc file`	count number of lines/words/characters in file

PREVIOUS

[Files in Unix](#)

NEXT

[More About Files](#)

Last updated on Apr 29, 2019

Search...

UNIX Tutorials for Beginners

Introduction to Unix

Tutorial 1: Files in Unix

Tutorial 2: Working with Files

Tutorial 3: More About Files

Tutorial 4: Wildcards and File Access

Tutorial 5: Processes and Jobs

Tutorial 6: Other Useful Commands

Tutorial 7: Environment Variables

# More About Files

## Redirection

Most processes initiated by Unix commands write to the *standard output* (that is, they write to the terminal screen), and many take their input from the *standard input* (that is, they read it from the keyboard). There is also the *standard error*, where processes write their error messages, by default to the terminal screen. Type `cat` without specifying a file to read

```
% cat
```

Then type a few words on the keyboard and press the [Return] key. Finally hold the [CTRL] key down and press `[d]` (written as `^D` for short) to end the input. What has happened? If you run the `cat` command without specifying a file to read, it reads the standard input (the keyboard), and on receiving the ‘end of file’ (`^D`), copies it to the standard output (the screen). In UNIX, we can redirect both the input and the output of commands.

## Redirecting the Output

We use the `>` symbol to redirect the output of a command. For example, to create a file called `list1` containing a list of fruit, type

```
% cat > list1
```

Then type in the names of some fruit as follows. Press [Return] after each one. Terminate with the end-of-file marker control-d (`^D`).

```
pear
banana
apple
^D
```

The `cat` command reads the standard input (the keyboard) and the `>` redirects the output, which normally goes to the screen, into a file called `list1`. To read the contents of the file, type

```
% cat list1
```

## Exercise 3A

Using the above method, create another file called `list2` containing the following fruit: orange, plum, mango, grapefruit. The form `>>` *appends* standard output to a file. So to add more items to the file `list1`, type

```
% cat >> list1
```

Then type in the names of more fruit

```
peach
grape
strawberry
^D
```

To read the contents of the file, type

```
% cat list1
```

You should now have two files. One contains six fruit names, the other contains four fruits. We will now use the `cat` command to join (concatenate) `list1` and `list2` into a new file called `biglist`. Type

```
% cat list1 list2 > biglist
```

This reads the contents of list1 and list2 in turn, then outputs the text to the file biglist. To read the contents of the new file, type

```
% cat biglist
```

## Redirecting the Input

We use the `<` symbol to redirect the input of a command. The command `sort` alphabetically or numerically sorts a list. Type

```
% sort
```

Using `<` you can redirect the input to come from a file rather than the keyboard. For example, to sort the list of fruit, type

```
% sort < biglist
```

and the sorted list will be output to the screen. To output the sorted list to a file, type

```
% sort < biglist > slist
```

Use `cat` to read the contents of the file `slist`.

## Pipes

To see who is on the system with you, type

```
% who
```

One method to get a sorted list of names is to type

```
% who > names.txt
% sort < names.txt
```

This is a bit slow and you have to remember to remove the temporary file called `names.txt` when you have finished. What you really want to do is connect the output of the `who` command directly to the input of the `sort` command. This is exactly what pipes do. The symbol for a pipe is the vertical bar `|` which, on a US keyboard, is above `Enter` on the right, with the backslash. For example, typing

```
% who | sort
```

will give the same result as above, but quicker and cleaner. To find out how many users are logged on, use `wc` (word count) with the option `-l` (`ell`) for number of lines only:

```
% who | wc -l
```

## Exercise 3B

Using pipes, print all lines of `list1` and `list2` containing the letter ‘`p`’, sort the result, and print to a file `sorted_plist.txt`. Hint: from `grep --help` find an option to print only the line, omitting the file name.

► Solution

## Summary

Command	Operation
<code>command &gt; file</code>	redirect standard output to a file
<code>command &gt;&gt; file</code>	append standard output to a file
<code>command &lt; file</code>	redirect standard input from a file

Command	Operation
command1	command2
cat file1 file2 > file0	concatenate file1 and file2 to file0
sort	sort data
who	list users currently logged in

PREVIOUS

[Working with Files](#)

NEXT

[Wildcards and File Access](#)

Last updated on Apr 29, 2019



© 2023 The Rector and Visitors of the University of Virginia



Search...

UNIX Tutorials for Beginners

Introduction to Unix

Tutorial 1: Files in Unix

Tutorial 2: Working with Files

Tutorial 3: More About Files

Tutorial 4: Wildcards and File Access

Tutorial 5: Processes and Jobs

Tutorial 6: Other Useful Commands

Tutorial 7: Environment Variables

# Wildcards and File Access

## Wildcards

### The Characters \* and ?

The character **\*** is called a wildcard, and will match against none or more character(s) in a file (or directory) name. For example, in your **unixstuff** directory, type

```
% ls list*
```

This will list all files in the current directory starting with list. Try typing

```
% ls *list
```

This will list all files in the current directory ending with list. The character **?** will match exactly one character. So **ls ?ouse** will match files like house and mouse, but not grouse. Try typing

```
% ls ?list
```

## Filename Conventions

Unix regards a directory as merely a special type of file. So the rules and conventions for naming files apply also to directories. In naming files, characters with special meanings, such as **/** **\*** **&** **%**, should be avoided. Also avoid using spaces within names. The safest way to name a file is to use only alphanumeric characters, that is, letters and numbers, together with **\_** (underscore) and **.** (dot). File names conventionally start with a lower-case letter, and may end with a dot followed by a group of letters indicating the contents of the file. For example, all files consisting of C code may be named with the ending **.c**, for example, **prog1.c**. Then in order to list all files containing C code in your home directory, you need only type **ls \*.c** in that directory.

Beware: some applications give the same name to all the output files they generate. For example, some compilers, unless given the appropriate option, produce compiled files named a.out. Should you forget to use that option, you are advised to rename the compiled file immediately, otherwise the next such file will overwrite it and it will be lost.

The Unix system itself ignores file suffixes. You could name your C++ program **mycode.py** and Linux will not care. However, many applications do care about file extensions. A C++ compiler will expect the file to end in **.cpp** or **.cxx**. It will not recognize a file ending in **.py**, but a Python interpreter would.

## Filesystem Security (Access Rights)

In your unixstuff directory, type

```
% ls -l (l for long listing!)
```

You will see that you now get lots of details about the contents of your directory, similar to the example below.



```
-rw----- 1 mst3k users      340 Jan 24 09:41 contour.py
-rw-r--r-- 1 mst3k users      322 May 29 2020 omparea.c
-rw-r--r-- 1 mst3k users      336 May 29 2020 omparea.f90
-rwxr-xr-x 1 mst3k users    105104 Jan 30 09:25 ompheatedplate
-rw-r--r-- 1 mst3k users     3144 Jan 24 11:49 ompheatedplate.cxx
-rw-r--r-- 1 mst3k users     2147 Jan 25 07:31 ompheatedplate.f90
-rw-r--r-- 1 mst3k users       302 Jan 30 13:39 omp.slurm
-rw-r--r-- 1 mst3k users 197321360 Jan 30 13:22 plate5
drwxr-sr-x 2 mst3k users     2560 May 29 2020 pythonMP
-rw-r--r-- 1 mst3k users       270 Jan 30 13:39 serial.slurm
-rw-r--r-- 1 mst3k users       403 Jan 30 15:52 slurm-46361608.out
```

Each file (and directory) has associated access rights, which may be found by typing `ls -l`.

In the left-hand column is a 10-symbol string consisting of the symbols `d`, `r`, `w`, `x`, `-`, and, occasionally, `s` or `S`. If `d` is present, it will be at the left hand end of the string and indicates a directory; otherwise `-` will be the starting symbol of the string. The nine remaining symbols indicate the permissions, or access rights, and are taken as three groups of three.

- the leftmost group of 3 gives the file permissions for the user that owns the file (or directory) (mst3k in the above example);
- the middle group gives the permissions for the group of people to whom the file (or directory) belongs (users in the above example);
- the rightmost group gives the permissions for all others. The symbols `r`, `w`, etc., have slightly different meanings depending on whether they refer to a simple file or to a directory.

Access Rights on Files

- `r` (or `-`), indicates read permission, that is, the presence or absence of permission to read and copy the file
- `w` (or `-`), indicates write permission, that is, the permission (or otherwise) to change a file
- `x` (or `-`), indicates execution permission, that is, the permission to execute a file, where appropriate

Access Rights on Directories

- `r` allows users to list files in the directory
- `w` means that users may delete files from the directory or move files into it
- `x` means the right to access files in the directory or to `cd` into it.

So, in order to read a file, you must have execute permission on the directory containing that file, and hence on any directory containing that directory as a subdirectory, and so on, up the tree.

Some Examples

Permissions	Meaning
<code>-rwxrwxrwx</code>	a file that everyone can read, write and execute (and delete)
<code>-rw-----</code>	a file that only the owner can read and write: no one else can read or write and no one has execution rights (e.g., yourmailbox file)

Changing Access Rights

`chmod` (changing a file mode)

Only the owner of a file can use `chmod` to change the permissions of a file. The options of `chmod` are as follows:

Symbol	Meaning
<code>u</code>	user
<code>g</code>	group

Symbol	Meaning
o	other
a	all
r	read
w	write (and delete)
x	execute (and access directory)
+	add permission
-	take away permission

For example, to remove read write and execute permissions on the file biglist for the group and others, type

```
% chmod go-rwx biglist
```

This will leave the other permissions unaffected. To give read and write permissions on the file biglist to all,

```
% chmod a+rw biglist
```

Exercise 4A

Try changing access permissions on the file science.txt and on the directory backups. Use ls -l to check that the permissions have changed.

Shared Systems

Some multiuser systems, such as high-performance clusters, may not permit chmod on certain directories, or may automatically revert to the default set of permissions (called the umask). This is for data protection and privacy for users.

Summary

Command	Operation
*	match any number of characters
?	match one character
chmod	change file or directory permissions

PREVIOUS

[More About Files](#)

NEXT

[Processes and Jobs](#)

Last updated on Apr 29, 2019

Search...

UNIX Tutorials for Beginners

Introduction to Unix

Tutorial 1: Files in Unix

Tutorial 2: Working with Files

Tutorial 3: More About Files

Tutorial 4: Wildcards and File Access

Tutorial 5: Processes and Jobs

Tutorial 6: Other Useful Commands

Tutorial 7: Environment Variables

# Processes and Jobs

A *process* is an executing program identified by a unique PID (process identifier). To see information about your processes, with their associated PID and status, type

```
% ps -u mst3k
```

A process may be in the foreground, in the background, or be suspended. In general the shell does not return the Unix prompt until the current process has finished executing. Some processes take a long time to run and hold up the terminal. Backgrounding a long process has the effect that the Unix prompt is returned immediately, and other tasks can be carried out while the original process continues executing.

## Running Background Processes

To background a process, type an & at the end of the command line. For example, the command sleep waits a given number of seconds before continuing. Type

```
% sleep 10
```

This will wait 10 seconds before returning the command prompt %. Until the command prompt is returned, you can do nothing in that terminal except wait. To run sleep in the background, type

```
% sleep 10 &

[1] 6259
```

The **&** runs the process in the background and returns the prompt immediately, allowing you to run other programs while waiting for that one to finish. The first line in the above example is typed in by the user; the next line, indicating *job number* and PID, is returned by the machine. The user is be notified of a job number (numbered from 1) enclosed in square brackets, together with a PID and is notified when a background process is finished. Backgrounding is useful for jobs which will take a long time to complete.

To the system, a “job” represents a group of processes (often just one) to which *signals* should be sent. Signals include backgrounding, foregrounding, and cancelling. This meaning of “job” must be distinguished from “jobs” submitted to a queueing system such as Slurm. To a resource manager (queueing system) a “job” is a set of instructions to be executed.

## Backgrounding a Current Foreground Process

At the prompt, type

```
% sleep 100
```

You can suspend the process running in the foreground by holding down the **[CTRL]** key and typing **[z]** (written as **^Z**) Then to put it in the background, type

```
% bg
```

Note: do not background programs that require user interaction.

## Listing Suspended and Background Processes

When a process is running, backgrounded or suspended, it will be entered onto a list along with a job number. To examine this list, type

```
% jobs
```

An example of a job list could be

1. Suspended sleep 100
2. Running firefox
3. Running vi

To restart (foreground) a suspended processes, type

```
% fg %jobnumber
```

For example, to restart sleep 100, type

```
% fg %1
```

Typing **fg** with no job number foregrounds the last suspended process.

## Killing a Process

### **kill** (terminate or signal a process)

It is sometimes necessary to kill a process (for example, when an executing program is in an infinite loop). To kill a job running in the foreground, type **^C** ([CTRL] + [c]). For example, run

```
% sleep 100 ^C
```

To kill a suspended or background process, type

```
% kill %jobnumber
```

For example, run

```
% sleep 100 &
```

```
% jobs
```

If it is job number 4, type

```
% kill %4
```

To check whether this has worked, examine the job list again to see if the process has been removed.

### **ps** (process status)

Alternatively, processes can be killed by finding their process numbers (PIDs) and using **kill PID\_number**:

```
% sleep 100 &
% ps

PID    TTY          TIME CMD
20077  pts/5  S        0:05 sleep 100
21563  pts/5  T        0:00 firefox
21873  pts/5  S        0:25 vi
```

To kill off the process **sleep 100**, type

```
% kill 20077
```

and then type **ps** again to see if it has been removed from the list. If a process refuses to be killed, uses the **-9** option, i.e., type

```
% kill -9 20077
```

Linux systems support a command **killall** which takes the name of the process rather than its PID.

```
killall -9 sleep
```

Note: It is not possible to kill off other users' processes! Unless, of course, it is a system you control and for which you have *root* privileges. The *root* account is the system administrator, and on Linux is all-powerful.

## Summary

Command	Operation
<code>ls -lag</code>	list access rights for all files
<code>chmod [options] file</code>	change access rights for named file
<code>command &amp;</code>	run command in background
<code>^C</code>	kill the job running in the foreground
<code>^Z</code>	suspend the job running in the foreground
<code>bg</code>	background the suspended job
<code>jobs</code>	list current jobs
<code>fg %1</code>	foreground job number <code>1</code>
<code>kill %1</code>	kill job number <code>1</code>
<code>ps</code>	list current processes
<code>kill 26152</code>	kill process number <code>26152</code>
<code>killall name</code>	kill process <code>name</code>

PREVIOUS

[Wildcards and File Access](#)

NEXT

[Other Useful Commands](#)

Last updated on Apr 29, 2019

Search...

UNIX Tutorials for Beginners

Introduction to Unix

Tutorial 1: Files in Unix

Tutorial 2: Working with Files

Tutorial 3: More About Files

Tutorial 4: Wildcards and File Access

Tutorial 5: Processes and Jobs

Tutorial 6: Other Useful Commands

Tutorial 7: Environment Variables

# Other Useful Commands

## exit

Exit the current shell. If it is the login shell, this command logs the user off.

## which

The **which** command indicates the path to the executable specified.

```
% which myexec
```

The **which** command returns the location of the executable according to the rules used to search paths. The shell always searches from left to right in the list contained in the PATH environment variable; the first executable of the specified name is the one that is used.

## wc

The **wc** command returns the number of lines, words, and characters in an ASCII file. A word is defined as a non-zero length string surrounded by whitespace.

```
% wc myfile.txt
```

To print only the number of lines, use

```
% wc -l myfile.txt
```

## diff

**diff** shows the differences between two ASCII files on a per-line basis.

```
% diff file1 file2
```

## find

**find** is an extremely powerful command with many options. The simplest and most common use of it is to search for a file of a given name or with a name containing a pattern.

```
% find . -name myscript.sh
```

This starts from current directory (.) and searches for myscript.sh. The period is optional under Linux (but not under Mac OSX). To search for a name with a pattern it must typically be enclosed in quotes

```
% find . -name "*.sh"
```

See the manpage or examples online for more usage patterns of this command.

## du

The **du** command outputs the number of kilobytes used by each subdirectory. Useful if you have gone over quota and you want to find out which directory has the most files. Some options make it more useful; in particular, -s summarizes directories and -h prints it in human-readable format. In your home directory, type

```
% du -s -h *
```

## gzip and zip

This reduces the size of a file, thus freeing valuable disk space. For example, type

```
% ls -l science.txt
```

and note the size of the file. Then to compress `science.txt`, type

```
% gzip science.txt
```

This will compress the file and place it in a file called `science.txt.gz`. To see the change in size, type `ls -l` again. To uncompress the file, use the `gunzip` command.

```
% gunzip science.txt.gz
```

Most Linux systems also provide the standard `zip` and `unzip` commands.

```
% zip science.txt.zip
% unzip science.txt.zip
```

## tar (tape archive)

The standard archive format in Linux is `tar`. Tar is usually used to bundle directories. (Zip can also be used for this purpose.) Typically the output file is compressed with `gzip`.

```
tar czf mydir.tar.gz mydir
```

The `c` option creates the tarfile (also called a “tarball”). The `f` option is for file, and this form of the command must be followed by the name of the file to contain the archive. The `z` option `gzi`ps the file.

To extract the files

```
tar xf mydir.tar.gz
```

Newer versions of `tar` can detect that the archive is zipped, so a `z` is not necessary for extraction. The `x` option extracts. This will create the directory `mydir` if it does not exist. If it does, the contents will be replaced by the contents of `mydir.tar.gz`.

## file

`file` classifies the named files according to the type of data they contain, for example ASCII (text), pictures, compressed data, etc. To report on all files in your home directory, type

```
% file *
```

## cut

The `cut` command extracts selected portions of a line, based on fields separated by a delimiter

```
% cut?d delim ?fC1,C2,C3
```

Examples:

```
% cut -d ' ' ?f1 /etc/resolve.conf
% cat myfile | cut -c 80
```

## sort

This command sorts lines of a text file, based on command-line options. The default is to sort alphabetically, based on lexicographical ordering (in which e.g. 100 comes before 2).

```
% sort mylist.txt
```

## uniq

Removes duplicate lines (file must be sorted first since it only compares lines pairwise).

```
% uniq mylist.txt
```

A frequent pattern is to pipe the output of sort into `uniq`

```
% sort animals | uniq
```

## history

The bash shell keeps an ordered list of all the commands that you have entered. Each command is given a number according to the order it was entered.

```
% history (show command history list)
```

If you are using the `bash` or `tcsh` shell, you can use the exclamation character (`!`) to recall commands easily.

```
% !! (recall last command)
% !-3 (recall third most recent command)
% !5 (recall 5th command in list)
% !grep (recall last command starting with grep)
```

You can increase the size of the history buffer by typing

```
% set history=100
```

### PREVIOUS

[Processes and Jobs](#)

### NEXT

[Environment Variables](#)

---

Last updated on Apr 29, 2019



Research Computing

© 2023 The Rector and Visitors of the University of Virginia





Search...

UNIX Tutorials for Beginners

Introduction to Unix

Tutorial 1: Files in Unix

Tutorial 2: Working with Files

Tutorial 3: More About Files

Tutorial 4: Wildcards and File Access

Tutorial 5: Processes and Jobs

Tutorial 6: Other Useful Commands

Tutorial 7: Environment Variables

# Environment Variables

Variables are a way of passing information from the shell to programs when you run them. Programs look “in the environment” for particular variables and if they are found will use the values stored. Some are set by the system, others by you, yet others by the shell, or any program that loads another program. Standard Unix variables are split into two categories, *environment variables* and *shell variables*. In broad terms, shell variables apply only to the current instance of the shell and are used to set short-term working conditions. Environment variables are exported and have a farther reaching significance; those set at login are valid for the duration of the session. By convention, environment variables are written in UPPERCASE while shell variables usually have lowercase names.

## Environment Variables

An example of an environment variable is the `$SHELL` variable. The value of this is the current shell you are using. Type

```
% echo $SHELL
```

More examples of environment variables are

```
$USER (your login name)
$HOME (the path name of your home directory)
$PWD (current working directory)
$DISPLAY (the name of the computer screen to display X windows; only set if X is enabled)
$PATH (the directories the shell should search to find a command)
```

## Using and Setting Variables

Environment variables are set using the `export` command (`bash` or `zsh`) or the `setenv` command (`tcsh` or `csh`), displayed using the `printenv` (`bash`, `tcsh`) or `env` (`bash`, `zsh`) commands, and unset using the `unsetenv` command. To show all values of these variables, type

```
% printenv | more
```

To set a value of an environment variable, type (for `bash`)

```
% export VAR=value
```

## Sourcing

A group of shell commands can be placed into a file and then *sourced*. When a file is sourced, the commands are executed as if they had been typed at the command line in the current shell. For example, if several environment variables needed to be set over and over again, they could be collected into a file such as this simple script called `envs.sh`:

```
export NUM_CPUS=16
export MEM=64

ncpus=${NUM_CPUS}
```



## Exercise 6A

Download the `envs.sh` file to the Unix system you are using. Run the command

```
source envs.sh
```

Print the values of the environment variables in the file. To print the value of the shell variable `ncpus`, type

```
echo $ncpus
```

## Dotfiles

Each time you log in to a Unix host, the system looks in your home directory for initialization files. Information in these files is used to set up your working environment. The first initialization file sourced is the *login* initialization. It is sourced only in the login shell. Note: modern “desktop” user interfaces tend to “swallow” the login setup file, and it may be difficult to determine what is happening in these cases if there is an error.

At login the bash shell first sources `.bash_profile` or `.profile` (if `.bash_profile` exists `.profile` will be ignored). Child shells source `.bashrc`. The `zsh` sources `.zprofile` and child shells source `.zshrc`. Two older shells, `csh` (C shell) and `tcsh`, read `.login` for login shells and `.cshrc` or `.tcshrc` for all other shells.

Note that all these file names begin with periods or “dots”; hence they are called *dotfiles*. As we have learned, dotfiles are hidden and will only be visible with `ls -a`.

The `.bash_profile`, `.profile`, or `.login` is to set conditions which will apply to the whole session and/or to perform actions that are relevant only at login. The `.bashrc`, `.zshrc`, or `.tcshrc` file is used to set conditions and perform actions specific to the shell and to each invocation of it. The `rc` stands for resource; many Unix dotfiles use this convention to set resources.

If you wish for your login shell to source the `.bashrc` also, add the lines

```
if [ -f ~/.bashrc ];  
then . ~/.bashrc  
fi
```

to the `.bash_profile` script.

Warning: NEVER put commands that run graphical displays (e.g. web browsers) in your dotfiles. If you change your `.bashrc` you can force the shell to reread it by using the shell `source` command.

```
% source ~/.bashrc
```

## Setting the Path

When you type a command, your path (or `$PATH`) variable defines in which directories the shell will look to find the command you typed. If the system returns a message saying “`command: Command not found`”, this indicates that either the command doesn’t exist at all on the system or it is simply not in your path.

For example, suppose you have installed a program called “units” into your home directory in a folder called `units174`. Units is a simple utility that can convert Imperial to metric and vice versa, from SI to cgi, and so forth. This folder contains a `bin` subdirectory in which the executable is located. To run units, you must either directly specify the units path (`~/units174/bin/units`), or you need to have the directory `~/units174/bin` in your path. You can add it to the end of your existing path (the `$PATH` represents this) by issuing the command:

```
% export PATH=$PATH:$HOME/units174/bin
```

If you have `units` installed you can test that this worked by trying to run units in any directory other than where units is actually located.

```
% cd; units
```

Hint: You can run multiple commands on one line by separating them with a semicolon.

To add this path permanently, add the `export` line to your `.bashrc` file after the list of other commands. Make sure that you include the `$PATH` when you reset it, or you will lose access to basic system commands!

#### PREVIOUS

[Other Useful Commands](#)

---

Last updated on Apr 29, 2019



Research Computing

© 2023 The Rector and Visitors of the University of Virginia

